



Bachelor's Thesis in Computer Science and Media

How does an Event Sourcing architecture compare to CRUD systems with an independent audit log, when it comes to scalability, performance and traceability?

Lukas Karsch

45259

Hochschule der Medien Stuttgart

Submitted on 2026/03/02

to obtain the degree of Bachelor of Science

Main Supervisor: Prof. Dr. Tobias Jordine

Secondary Supervisor: Felix Messner

Ehrenwörtliche Erklärung

Contents

1	Introduction	3
1.1	Motivation	3
1.2	Research question(s)	3
1.3	Goals and non goals	3
1.4	Structure of the paper	3
2	Basics	3
2.1	WWW, Web APIs, REST	3
2.2	Layered Architecture Foundations	4
2.3	Domain Driven Design	5
2.4	CRUD architecture	7
2.5	CQRS Architecture	7
2.6	(Eventual) Consistency	8
2.7	Event Sourcing and event-driven architectures	8
2.8	Traceability and auditing in IT systems	10
2.8.1	Why is traceability a business requirement	10
2.8.2	Audit Logs	10
2.8.3	Event Streams	10
2.8.4	Rebuilding state from an audit log and an event stream	10
2.9	Scalability of systems	10
3	Related Work	10
4	Proposed Method	10
4.1	Project requirements	10
4.2	Performance	10
4.3	Scalability or flexibility (TODO)	10
4.4	Traceability	10
4.5	Tech Stack	10
5	Implementation	10
5.1	CRUD implementation	10
5.2	ES/CQRS implementation	10
5.3	Infrastructure	10
6	Results	10
7	Discussion	10
7.1	Analysis of results	10
7.2	Conclusion & Further work	10

1 Introduction

1.1 Motivation

1.2 Research question(s)

1.3 Goals and non goals

1.4 Structure of the paper

2 Basics

2.1 WWW, Web APIs, REST

The World Wide Web (WWW) is a connected information network used to exchange data. Resources are can be accessed via Uniform Resource Identifiers (URIs) which are transferred using formats like JSON or HTML via protocols like HTTP. HTTP is a stateless protocol based on a request-response structure. It supports standardized request types, such as GET and POST, which convey a semantic meaning (Jacobs & Walsh 2004).

Web APIs are interfaces that enable applications to communicate. They use HTTP as a network-based API (Fielding 2000, p. 138). Modern APIs typically follow REST principles. REST stands for "Representational State Transfer" and describes an architectural style for distributed hypermedia systems (Fielding 2000, p. 76).

REST APIs adhere to principles derived from a set of constraints imposed by the HTTP protocol, for example. One such constraint is "stateless communication": Communication between clients and the server must be *stateless*, meaning the client must provide all the necessary information for the server to fully understand the request.

Furthermore, every resource in REST applications must be addressable via a unique ID, which can then be used to derive a URI to access the resource. Below are some examples for resources and URIs which could be derived from them:

- Book; ID=1; URI=`http://example.com/books/1`
- Book; ID=2; URI=`http://example.com/books/2`
- Author; ID=100; URI=`http://example.com/authors/100`

The "Hypermedia as the engine of application state (HATEOAS)" principle states that resources should be linked to each other. Clients should be

able to control the application by following a series of links provided by the server (Tilkov 2007).

Every resource must support the same interface, usually HTTP methods (GET, POST, PUT, etc.) where operations on the resource correspond to one method of the interface. For example, a POST operation on a customer might map to the `createCustomer()` operation on a service.

Resources are decoupled from their representations. Clients can request different representations of a resource, depending on their needs (Tilkov 2007): a web browser might request HTML, while another server or application might request XML or JSON.

2.2 Layered Architecture Foundations

Layered Architecture is the most common architecture pattern in enterprise applications. Applications following a layered architecture are divided into *horizontal layers*, with each layer performing a specific role. A standard implementation consists of the following layers:

- Presentation: Handles requests and displays data in a user interface or by turning it into representations (e.g. JSON)
- Business: Encapsulates business logic
- Persistence: Persists data by interacting with the underlying persistence technologies (e.g. SQL databases)
- Database

A key concept in this design is layers of isolation, where layers are "closed", meaning a request must pass through the layer directly below it to reach the next, ensuring that changes in one layer do not affect others.

In a layered application, data flows downwards during request handling and upwards during the response: a request arrives in the presentation layer, which delegates to the business layer. The business layer fetches data from the persistence layer which holds logic to retrieve data, e.g. by encapsulating SQL statements.

The database responds with raw data, which is turned into a Data Access Object by the persistence layer. The business layer uses this data to execute rules and make decisions. The result will be returned to the presentation layer which can then wrap the response and return it to the caller. (Richards 2015)

The data in layered applications is often times modeled in an *anemic* way. In an Anemic Domain Model, business entities are treated as only data. They are objects which contain no business logic, only getters and setters. Business logic is entirely contained in the business (or "service") layer. Fowler (2003) describes this as an object-oriented *antipattern*.

2.3 Domain Driven Design

Domain Driven Design (DDD) is a different architectural approach for applications. It differs from layered architecture primarily in the way the domain is modelled and the responsibilities of application services.

The core idea of DDD is that the primary focus of a software project should not be the underlying technologies, but the domain. The domain is the topic with which a software concerns itself. The software design should be based on a model that closely matches the domain and reflects a deep understanding of business requirements. (Evans 2004, pp. 8, 12)

This domain model is built from a *ubiquitous language* which is a language shared between domain experts and software experts. This ubiquitous language is built directly from the real domain and must be used in all communications regarding the software. (Evans 2004, pp. 24–26)

The software must always reflect the way that the domain is talked about. Changes to the domain and the ubiquitous language must result in an immediate change to the domain model.

When modeling the domain model, the aim should not be to create a perfect replica of the real world. While it should carefully be chosen, the domain model artificial and forms a selective abstraction which should be chosen for its utility. (Evans 2004, pp. 12, 13)

While Layered Architecture organizes code into technical tiers and is typically built on Anemic Domain Models, often resulting in the *big ball of mud* antipattern (Richards 2015, p. V), DDD demands a Rich Domain Model where objects incorporate both data and the behavior or rules that govern that data. The code is structured semantically into bounded context and modules which are chosen to tell the "story" of a system rather than its technicalities. (Evans 2004, p. 80)

Entities (also known as reference objects) are domain elements fundamentally defined by a thread of continuity and identity rather than their specific attributes. Entities must be distinguishable from other entities, even if they share the same characteristics. To ensure consistency and identity, a unique identifier is assigned to entities. This identifier is immutable throughout the object's life. (Evans 2004, pp. 65–69)

Value Objects are elements that describe the nature or state of something and have no conceptual identity of their own. They are interesting only for their characteristics. While two entities with the same characteristics are considered as different from each other, the system does not care about "identity" of a value object, since only its characteristics are relevant. Value objects should be used to encapsulate concepts, such as using an "Address" object instead of distinct "Street" and "City" attributes. Value objects should be immutable. They are never modified, instead they are replaced entirely when a new value is required. (Evans 2004, pp. 70–72)

Using a Rich Domain Model does not mean that there should be no layers, the opposite is the case. Evans (2004) advocates for using layers in domain driven designs. He proposes the following layers: (Evans 2004, p. 53)

- Presentation: Presents information and handles commands
- Application Layer: Coordinates app activity. Does not hold business logic, but delegate tasks and hold information about their progress
- Domain Layer: Holds information about the domain. Stateful objects (rich domain model) that hold business logic and rules
- Infrastructure layer: Supports other layers. Handles concerns like communication and persistence

Evans (2004, p. 75) points out that in some cases, operations in the domain can not be mapped to one object. For example, transferring money does conceptually not belong to one bank account. In those cases, where operations are important domain concepts, domain services can be introduced as part of model-driven design. To keep the domain model rich and not fall back into procedural style programming like with Anemic Domain Model, it is important to use services only when necessary. Services are not allowed to strip the entities and value objects in the domain of behavior. According to Evans, a good domain service has the following characteristics:

- The operation relates to a domain concept which would be misplaced on an entity or a value object
- The operation performed refers to other objects in the domain
- The operation is stateless

2.4 CRUD architecture

Layered architectures are the standard for data-oriented enterprise applications. These applications mostly follow a CRUD architecture. CRUD is an acronym coined by Martin (1983) that stands for "Create, Read, Update, Delete". These four actions can be applied to any record of data.

The state of domain objects in a CRUD architecture is often mapped to normalized tables on a relational database, though other storage mechanisms maybe used. The application acts on the current state of the data, with all actions (reads and writes) acting on the same data.

ACID (Atomicity, Consistency, Isolation, Durability) are an important feature of CRUD applications. They can be guaranteed using transactions, ensuring that data stays consistent and operations are atomic. (Bernstein & Newcomer 2009, pp. 10, 11)

2.5 CQRS Architecture

Command Query Responsibility Segregation (CQRS) is an architectural pattern based on the fundamental idea that the models used to update information should be separate from the models used to read information. This approach originated as an extension of Bertrand Meyer's Command And Query Separation (CQS) principle, which states that a method should either perform an action (a command) or return data (a query), but never both. (Meyer 2006, p. 148)

CQRS is different from CQS in the fact that in CQRS objects are split into two objects, one containing commands, one containing queries. (Young 2010, p. 17)

CQRS applications are typically structured by splitting the application into two paths:

- Command Side: Deals with data changes and captures user intent. Commands tell the system *what* needs to be done rather than over-writing previous state. Commands are validated by the system before execution and can be rejected. (Young 2010, pp. 11, 12)
- Read Side: Strictly for reading data. The read side is not allowed to modify anything in the data store. The read side typically stores Data Transfer Objects (DTOs) that can directly be returned to the presentation layer. (Young 2010, p. 20)

In a CQRS architecture, the read side typically updates its data asynchronously by consuming notifications or events generated by the write side.

Because the models for updating and reading information are strictly separated, a synchronization mechanism is required to ensure the read store eventually reflects the changes made by commands. This usually leads to stale data on the read side.

Each read service independently updates its model by consuming notifications or events published by the write side, allowing the read model to store optimized, denormalized views on the data. (Young 2010, p. 23)

2.6 (Eventual) Consistency

Gray et al. (1996) explain that large-scale systems become unstable if they are held consistent at all times according to ACID principles. This is mostly due to the large amount of communication necessary to handle atomic transactions in distributed systems. To address these issues, modern distributed systems often adopt the BASE (Basically Available, Soft State, Eventual Consistency) model which explicitly trades off isolation and strong consistency for availability. Eventually consistent systems are allowed to exist in a so-called "soft state" which eventually converges through the use of synchronization mechanisms over time rather than being strongly consistent at all times. (Braun et al. 2021; Vogels 2009) This creates an inconsistency window in which data is not consistent across the system. During this window, stale data may be read. (Vogels 2009)

2.7 Event Sourcing and event-driven architectures

Event driven architecture is a design paradigm where systems communicate via the production and consumption of events. Events are records of changes in the system's domain. (Michelson 2006) This approach allows for a high degree of loose coupling, as the system publishing an event does not need to know about the recipient(s) or how they will react. These architectures offer excellent horizontal scalability and resilience, as individual system components can fail or be updated without bringing down the entire network. (Fowler 2005)

Event Sourcing is an architectural pattern within the landscape of event driven architectures. Event-sourced systems ensure that *all* changes to a system's state are captured and stored as an ordered sequence of domain events. (Fowler 2005) Unlike traditional persistence models that overwrite data and store only the most recent state, event sourcing maintains an immutable record of every action taken over time. These events are persisted in an append-only event store, which serves as the principal source of truth

from which the current system state can be derived. (Fowler 2017; Lima et al. 2021)

The current state of any entity in such a system can be rebuilt by replaying the history of events from the log, starting from an initial blank state. (Fowler 2005) To address the performance costs of replaying thousands of events for every request, developers implement projections or materialized views, which are read-only, often denormalized versions of the data optimized for specific queries. (Malyi & Serdyuk 2024) This separation of concerns is frequently managed by pairing event sourcing with the Command Query Responsibility Segregation (CQRS) pattern, which physically divides the data structures used for reading from those used for writing state changes. (Young 2010, p. 50)

Because every action taken on the system is stored, a number of facilities can be built on top of the event log: Temporal queries can be made, which determine the exact state of the application at any point in time. The event log acts as an immutable audit trail, making Event Sourcing architectures highly valuable for systems like accounting applications. (Fowler 2005)

- 2.8 Traceability and auditing in IT systems**
 - 2.8.1 Why is traceability a business requirement**
 - 2.8.2 Audit Logs**
 - 2.8.3 Event Streams**
 - 2.8.4 Rebuilding state from an audit log and an event stream**
- 2.9 Scalability of systems**

3 Related Work

4 Proposed Method

- 4.1 Project requirements**
- 4.2 Performance**
- 4.3 Scalability or flexibility (TODO)**
- 4.4 Traceability**
- 4.5 Tech Stack**

5 Implementation

- 5.1 CRUD implementation**
- 5.2 ES/CQRS implementation**
- 5.3 Infrastructure**

6 Results

7 Discussion

- 7.1 Analysis of results**
- 7.2 Conclusion & Further work**

Finally, I'm done!

Glossary

ACID Atomicity, Consistency, Isolation, Durability. 7, 8

Anemic Domain Model The objects describing the domain only hold data, no logic. 5, 6

Atomicity Atomicity means that an action is either fully executed or not at all. Atomic operations make sure the application is not left in an invalid state (Bernstein & Newcomer 2009, p. 10). 7

BASE Basically Available, Soft State, Eventual Consistency. 8

CQRS Command Query Responsibility Segregation. 7, 9

CQS Command And Query Separation. 7

CRUD Create Read Update Delete. 7

DAO Data Access Object. 4

DDD Domain Driven Design. 5

DTO Data Transfer Object. 7

HATEOAS Hypermedia as the engine of application state. 3

HTML HyperText Markup Language. 4

HTTP HTTP stands for *Hypertext Transfer Protocol*. It is a protocol used in internet communication and was defined in RFC 2616 (*RFC 2616: HTTP/1.1* 2025). 3

JSON JavaScript Object Notation. 4

REST REST stands for *Representational State Transfer*. It is an architectural style for distributed hypermedia systems. 3

Rich Domain Model Objects incorporate both data and the behavior or rules that govern that data. 5, 6

URI Uniform Resource Identifier. 3

WWW World Wide Web. 3

XML Extensible Markup Language. 4

References

- Bernstein, Philip A. & Eric Newcomer (2009). *Principles of transaction processing*. eng. 2nd ed. The Morgan Kaufmann series in data management systems. Burlington, MA: Morgan Kaufmann Publishers. ISBN: 978-1-55860-623-4.
- Braun, Susanne et al. (Oct. 2021). “Tackling Consistency-related Design Challenges of Distributed Data-Intensive Systems - An Action Research Study”. In: *Proceedings of the 15th ACM / IEEE International Symposium on Empirical Software Engineering and Measurement (ESEM)*. arXiv:2108.03758 [cs], pp. 1–11. DOI: 10.1145/3475716.3475771. URL: <http://arxiv.org/abs/2108.03758> (visited on 12/27/2025).
- Evans, Eric (2004). *Domain-driven design: tackling complexity in the heart of software*. eng. Boston: Addison-Wesley. ISBN: 978-0-321-12521-7.
- Fielding, Roy Thomas (2000). “Architectural Styles and the Design of Network-based Software Architectures”. en. In.
- Fowler, Martin (Nov. 2003). *Anemic Domain Model*. URL: <https://martinfowler.com/bliki/AnemicDomainModel.html> (visited on 12/27/2025).
- (Dec. 2005). *Event Sourcing*. URL: <https://martinfowler.com/eaaDev/EventSourcing.html> (visited on 11/13/2025).
- (Feb. 2017). *What do you mean by “Event-Driven”?* URL: <https://martinfowler.com/articles/201701-event-driven.html> (visited on 11/13/2025).
- Gray, Jim et al. (1996). “The dangers of replication and a solution”. en. In: *Proceedings of the 1996 ACM SIGMOD international conference on Management of data*. DOI: 10.1145/233269.233330. URL: <https://dl.acm.org/doi/epdf/10.1145/233269.233330> (visited on 12/28/2025).
- Jacobs, Ian & Norman Walsh (Dec. 2004). *Architecture of the World Wide Web, Volume One*. URL: <https://www.w3.org/TR/webarch/> (visited on 12/27/2025).
- Lima, Stanley et al. (June 2021). “Improving observability in Event Sourcing systems”. In: *Journal of Systems and Software* 181, p. 111015. DOI: 10.1016/j.jss.2021.111015.
- Malyi, Roman & Pavlo Serdyuk (Aug. 2024). “Developing a Performance Evaluation Benchmark for Event Sourcing Databases”. en. In: *Vіsnik Nacional'nogo universitetu "L'viv's'ka politehnika". Seriâ Ìnformacijni sistemi ta mereži* 15, pp. 159–168. ISSN: 2524065X, 26630001. DOI: 10.23939/sisn2024.15.159. URL: <https://science.lpu.ua/sisn/all-volumes-and-issues/volume-15-2024/developing-performance-evaluation-benchmark-event> (visited on 11/03/2025).

- Martin, James (1983). *Managing the data-base environment*. eng. Englewood Cliffs, N.J.: Prentice-Hall. ISBN: 978-0-13-550582-3.
- Meyer, Bertrand (2006). “STANDARD EIFFEL”. en. In: 3.
- Michelson, Brenda (Feb. 2006). *Event-Driven Architecture Overview*. en. Tech. rep. 681. Boston, MA: Patricia Seybold Group, p. 681. DOI: 10.1571/bda2-2-06cc. URL: <http://www.customers.com/articles/event-driven-architecture-overview> (visited on 01/02/2026).
- RFC 2616: HTTP/1.1* (2025). URL: <https://www.w3.org/Protocols/HTTP/1.1/rfc2616.pdf> (visited on 12/27/2025).
- Richards, Mark (2015). “Software Architecture Patterns”. en. In.
- Tilkov, Stefan (2007). “A Brief Introduction to REST”. en. In.
- Vogels, Werner (Jan. 2009). “Eventually consistent”. en. In: *Communications of the ACM* 52.1, pp. 40–44. DOI: 10.1145/1435417.1435432. URL: <https://dl.acm.org/doi/epdf/10.1145/1435417.1435432> (visited on 12/28/2025).
- Young, Greg (2010). “CQRS Documents by Greg Young”. en. In: URL: https://cqrss.wordpress.com/wp-content/uploads/2010/11/cqrs_documents.pdf (visited on 12/26/2025).